



Zero ETL on AWS

Accelerating Insights, Eliminating Complexity





Are you a **data engineer**
or **data analyst**?



Do you worry about **data pipelines** or analytics
on operational data?

Agenda

Data pipelines: overview and challenges

AWS helps simplify data pipelines

AWS Zero-ETL features

Demo


Resources

Speaker



Sanchit Jain

FOLLOW ME



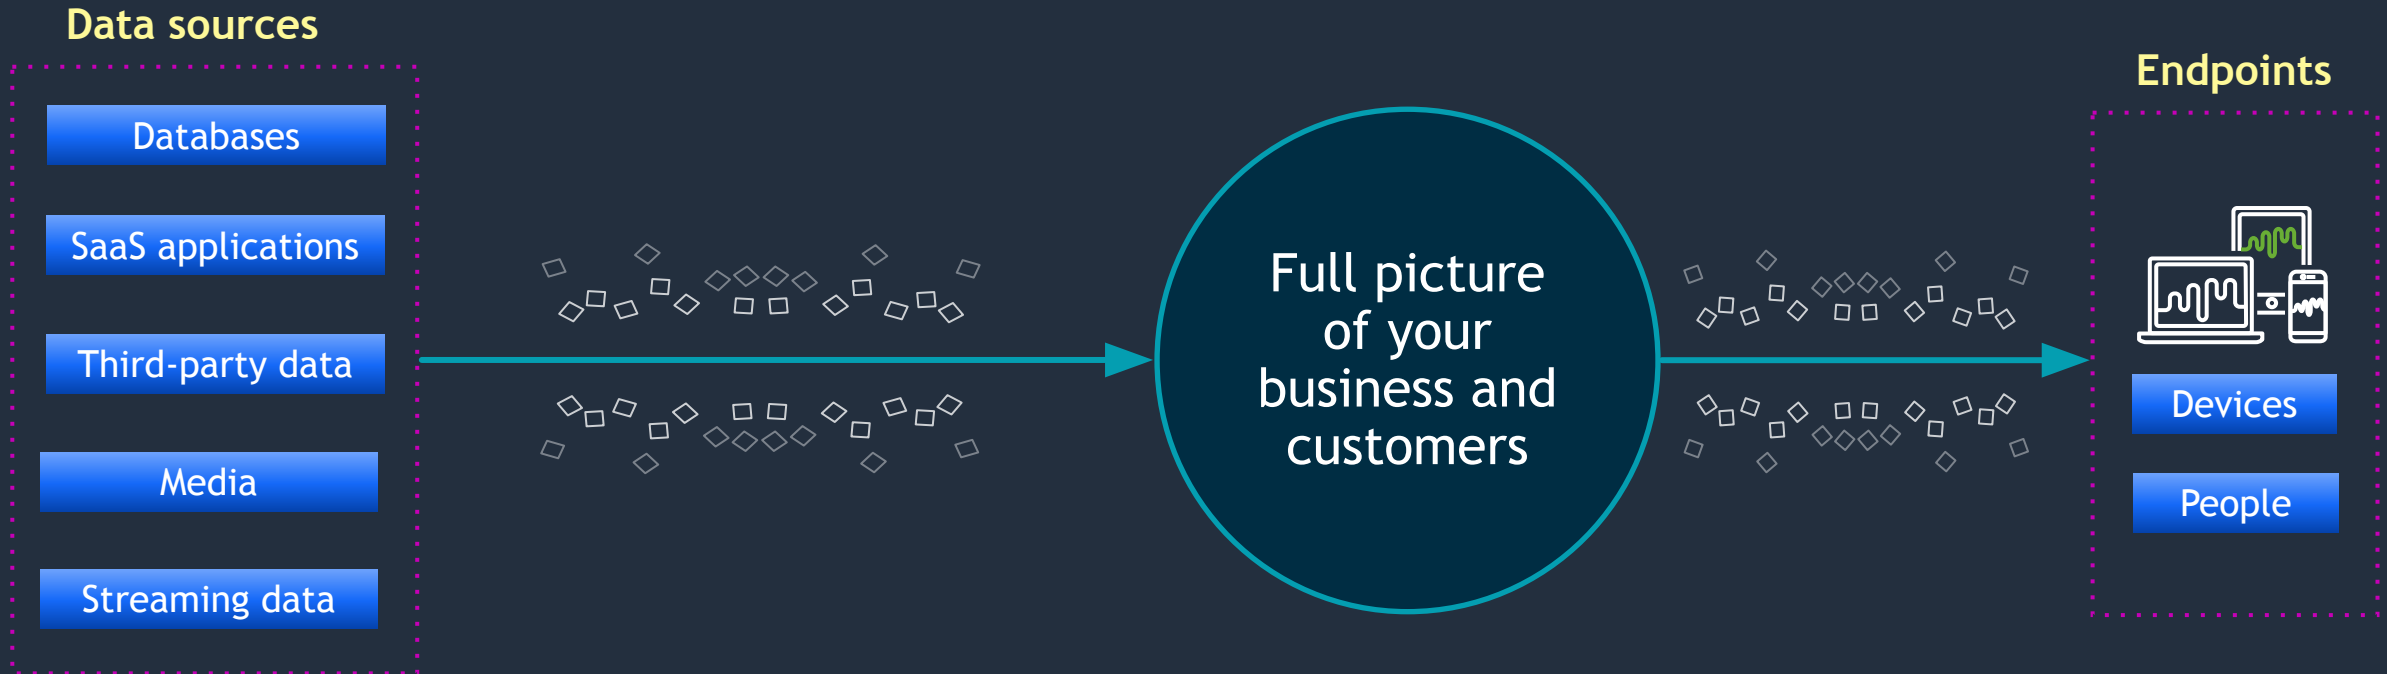
- Working as Data Analytics & Cloud Practice Lead - AWS Alliance at Quantiphi
- Recognition - AWS Hero & AWS Ambassador, MVP by Multicloud4u & AWS User Group Mumbai Lead
- Speaker at various events like AWS re:Invent, AWS Summit, AWS User Groups, AWS Community Days, and various educational institutes
- Follow me @ LinkedIn, Twitter or Github



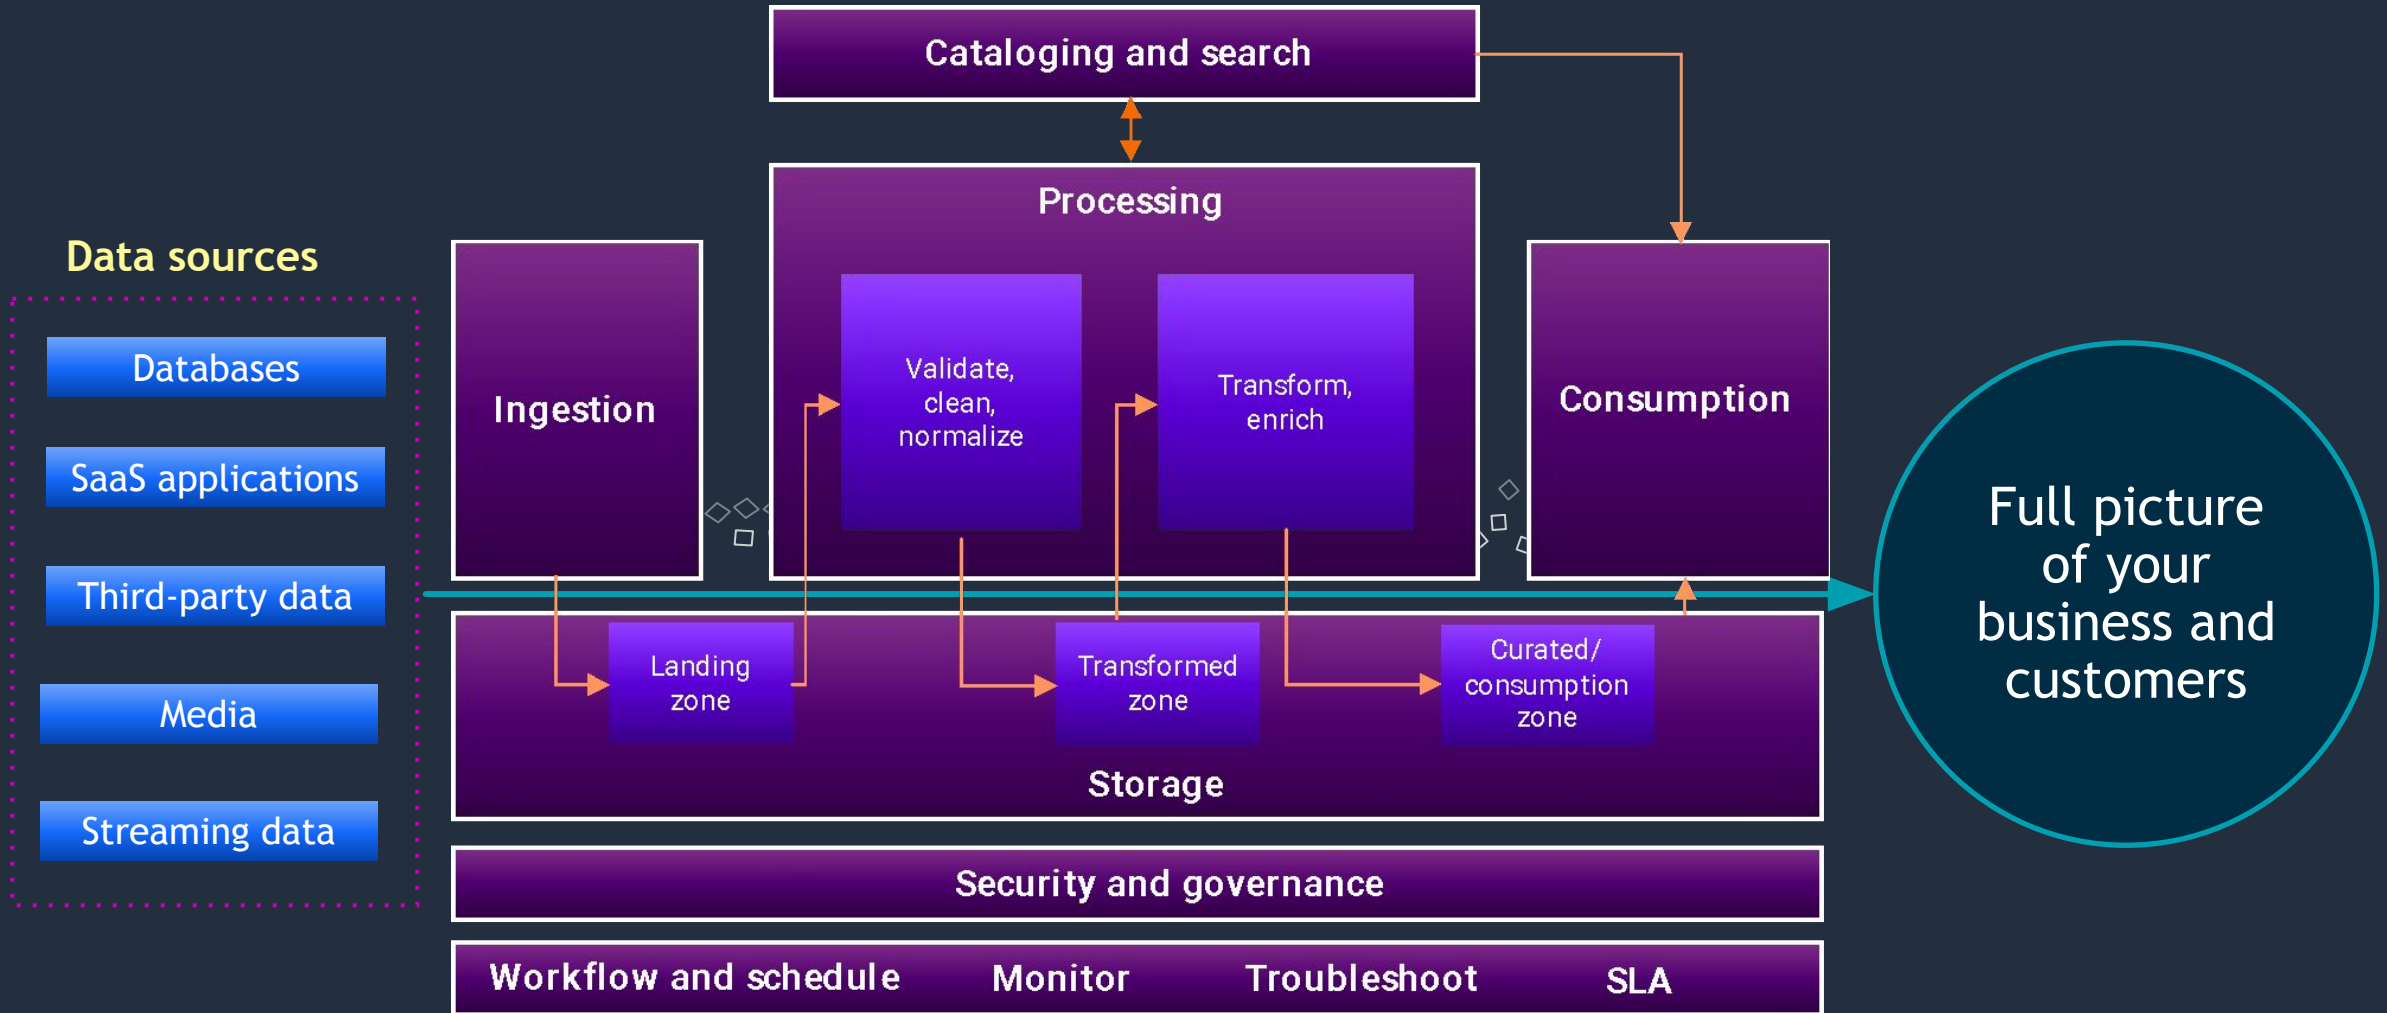
Data pipelines: overview and challenges



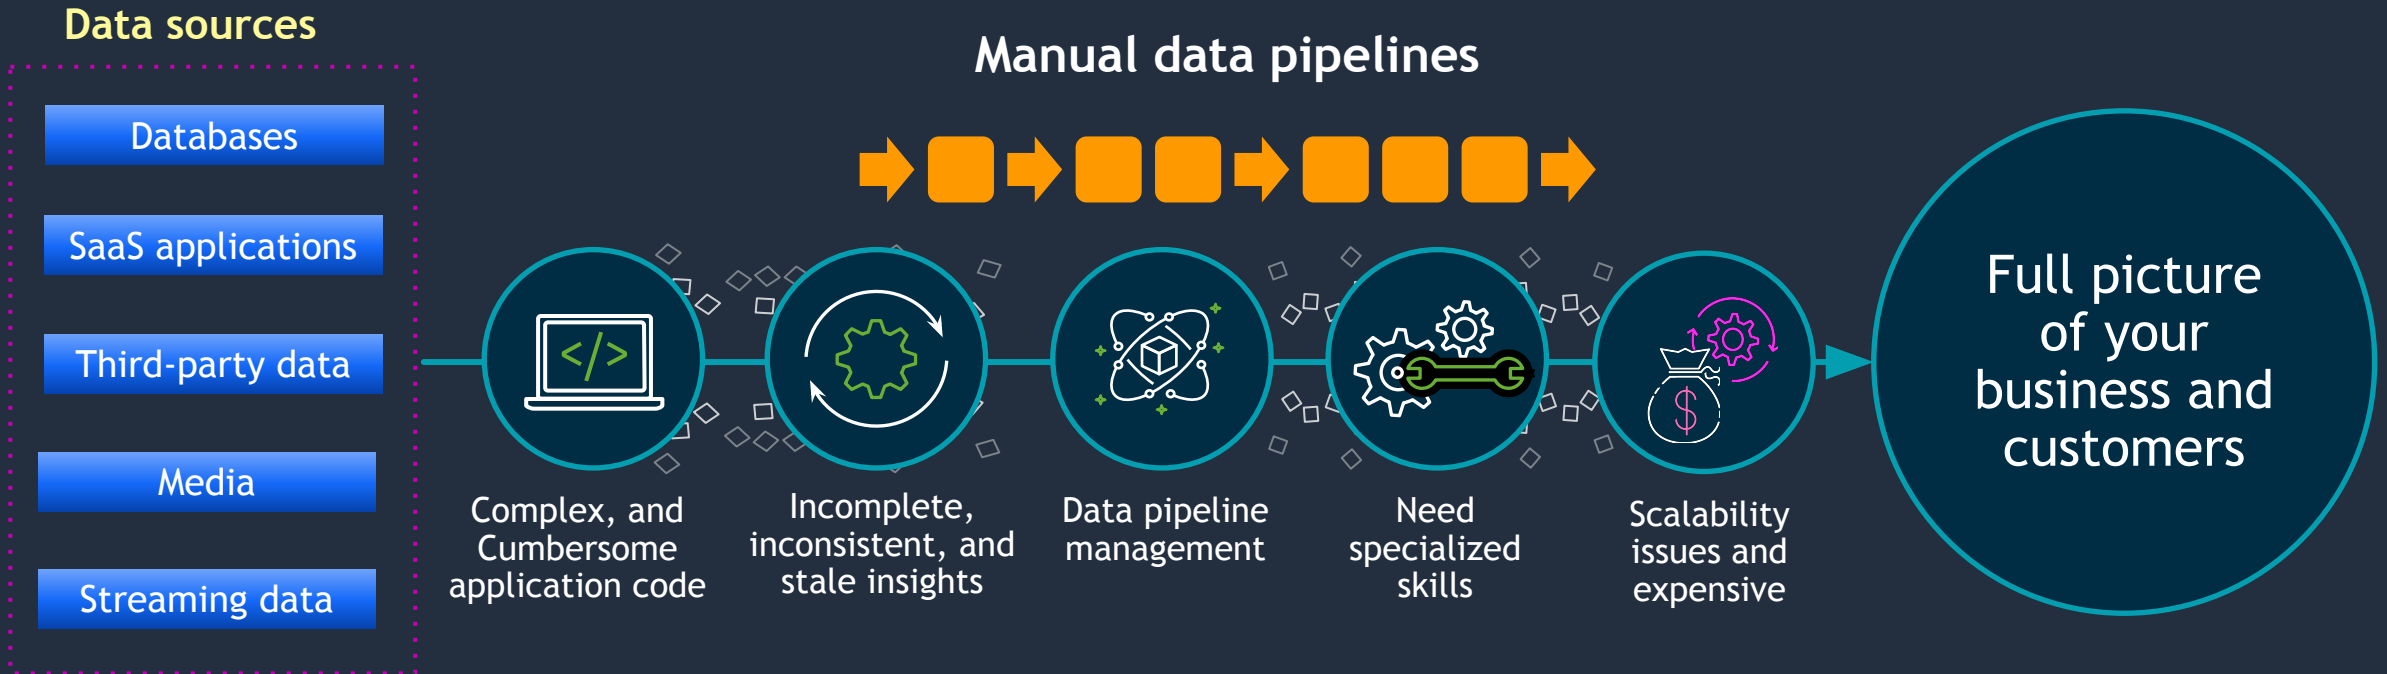
Data integration gives you a full picture of your business



Data integration and processing overview



Data integration and processing challenges






AWS helps simplify data
pipelines



Easily connect to all your data



Services for
**connecting to 100s
of data sources**




AWS Glue for
value-add data
transformations and more



Zero-ETL
future enabled
by service integrations


Services for connecting to 100s of data sources

50+




Amazon Appflow

30+




Amazon Kinesis Data Firehose

25+




Amazon Athena

40+




Amazon SageMaker

30+




Amazon QuickSight




AWS DataSync

300+



AWS Data Exchange

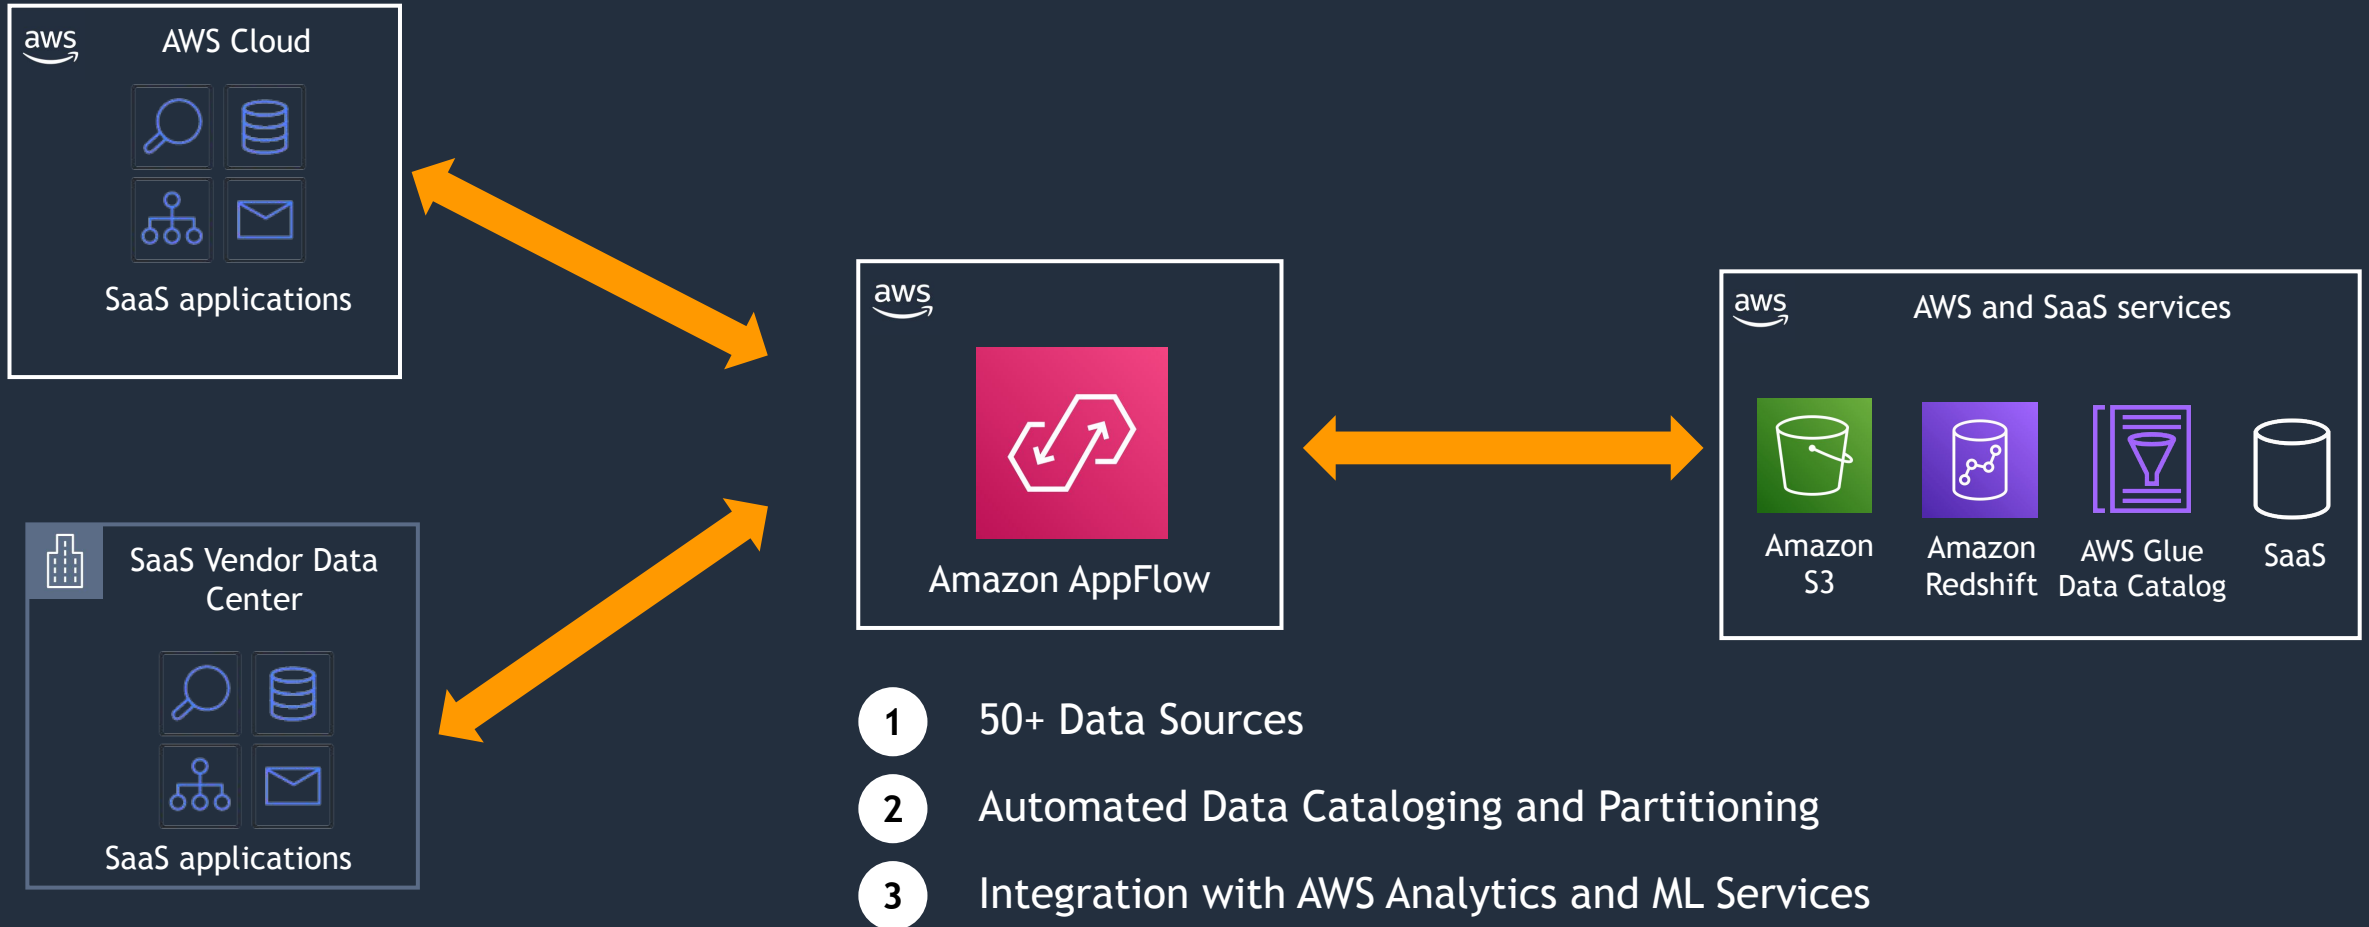
70+



AWS Glue

Integrating SaaS applications data

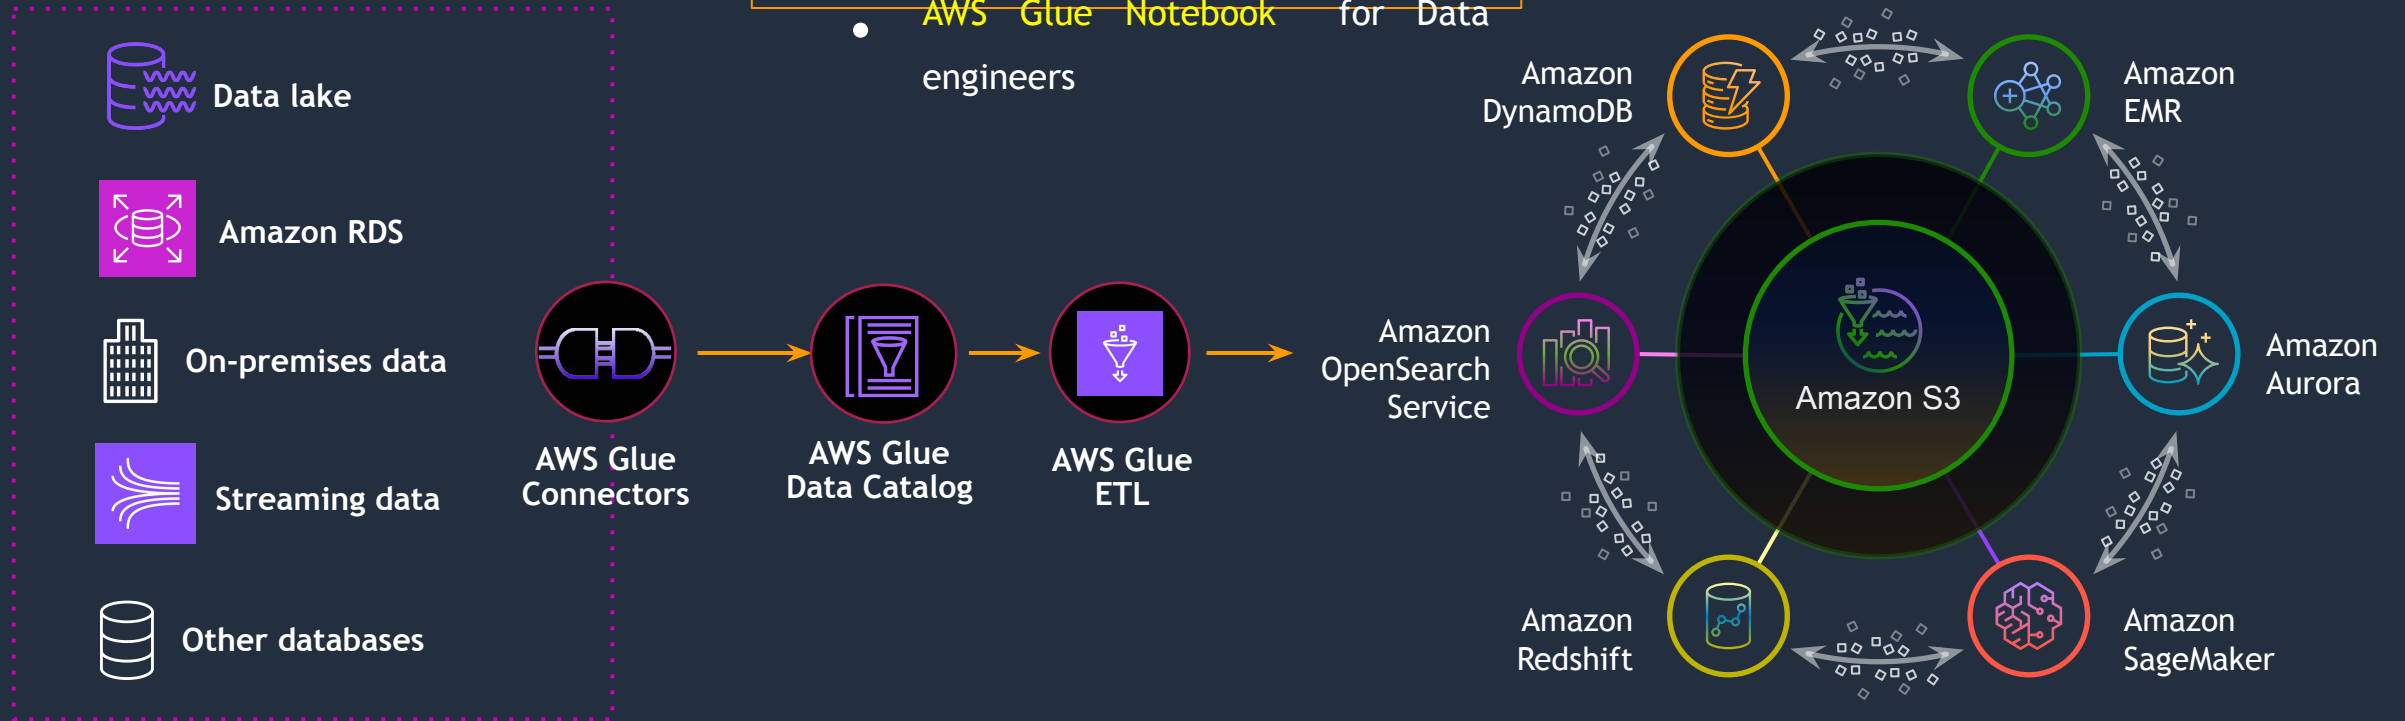
Amazon AppFlow makes it easy to ingest SaaS applications data into storage layer



Seamless data integration with AWS Glue

Persona-specific tools

- **AWS Glue Studio** for ETL developers
- **AWS Glue DataBrew** for Business analysts
- **AWS Glue Notebook** for Data engineers



Easily connect to all your data



Services for
**connecting to 100s
of data sources**

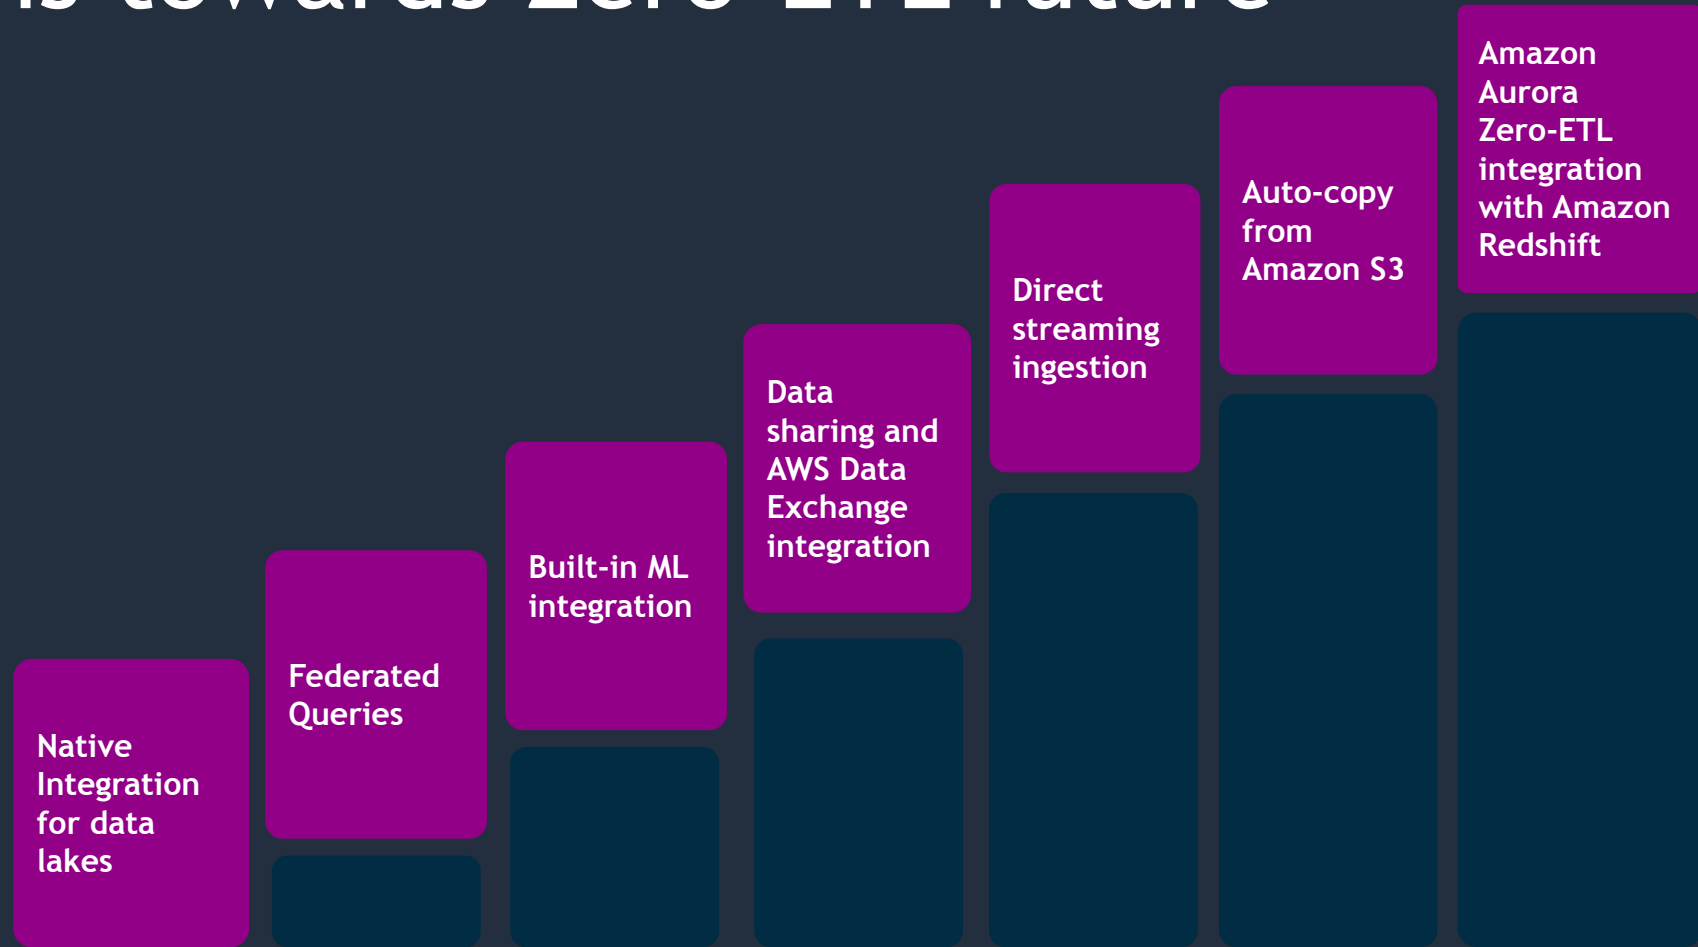


AWS Glue for
value-add data
transformations and more



Zero-ETL
future enabled
by service integrations

Innovations towards Zero-ETL future



ETL

Zero-ETL □

Unified analysis: Amazon Athena federated queries

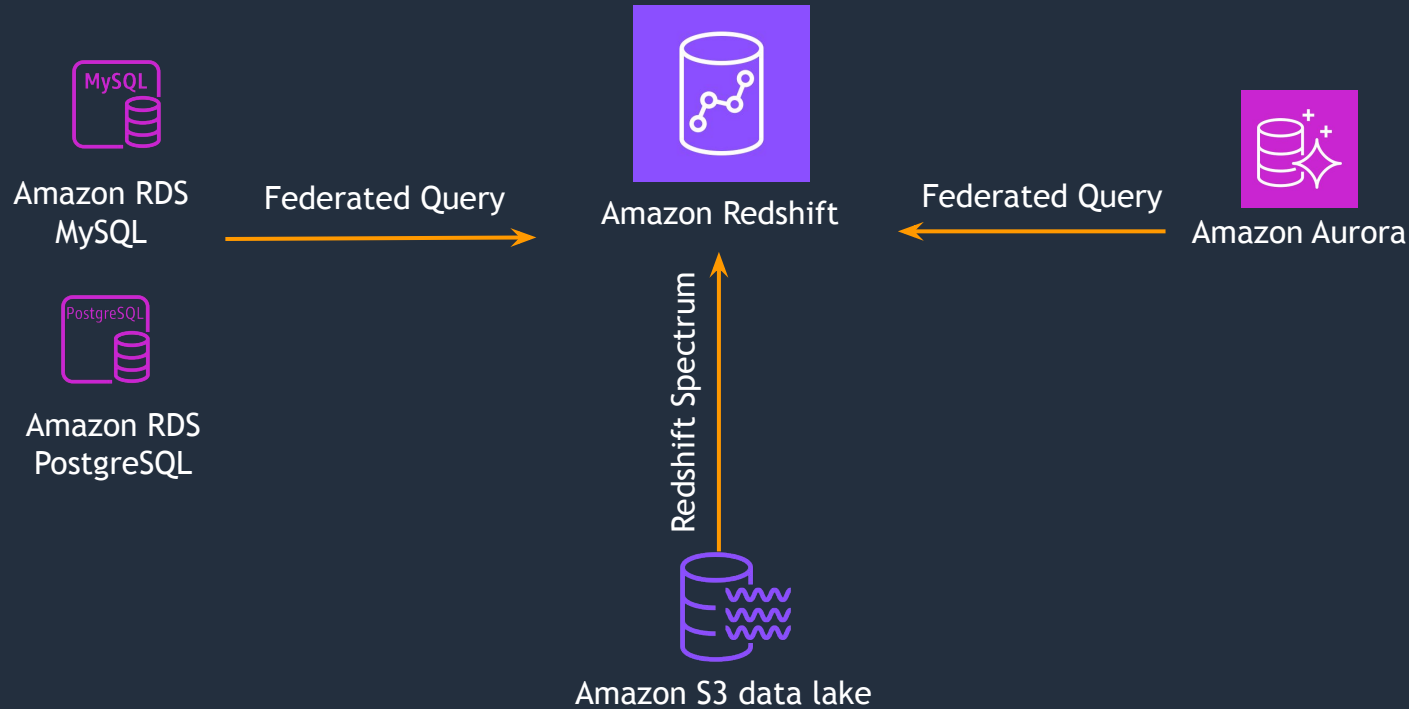


Run query across relational, non-relational, object, or custom data sources

Integrate external sources with your data lake without ETL

Run query across on-premises or cloud data sources

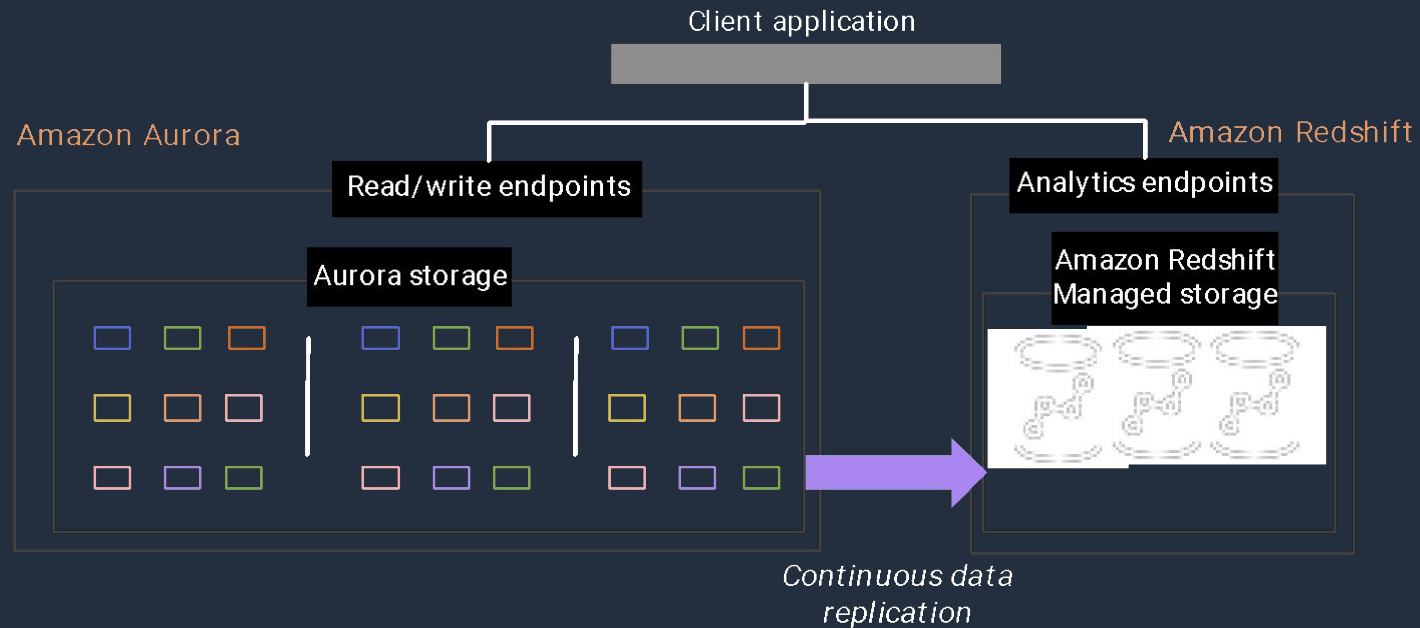
Redshift federated queries and Redshift Spectrum



- **Integrate operational databases** (Amazon RDS, Amazon Aurora) with data warehouse and data lake for real-time analytics
- Analytics on operational data **without data movement** and ETL delays
- **Flexible and easy** way to ingest data, avoiding complex ETL pipelines
- Support for open table formats like Apache Hudi and Delta Lake

Amazon Aurora zero-ETL integration with Amazon Redshift

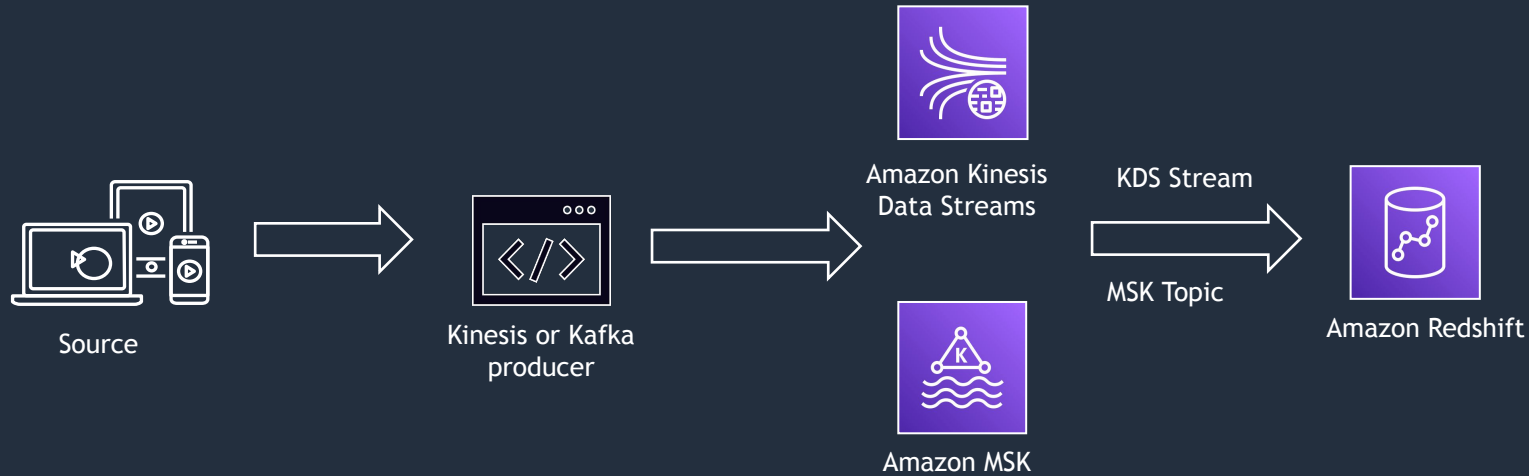
BOOST ANALYTICAL QUERY PERFORMANCE



- No ETL pipelines
- Real-time analytics and machine learning on transactional data
- Consolidated insights from multiple Amazon Aurora databases

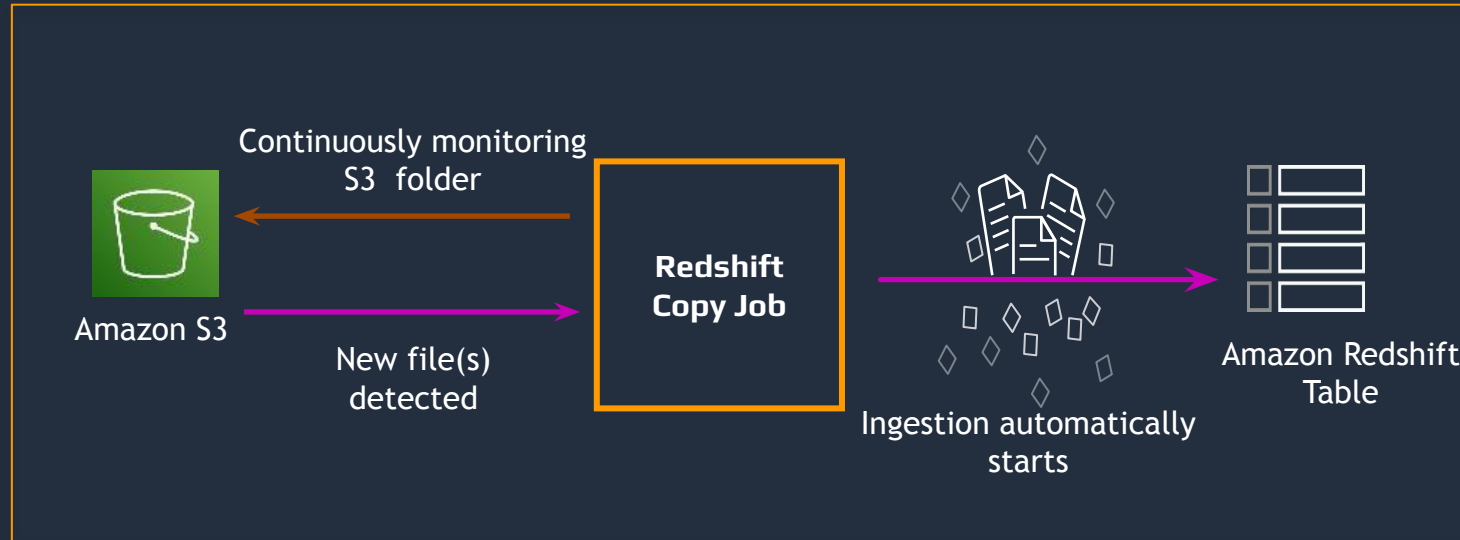
Amazon Redshift Streaming Ingestion support

NATIVE INTEGRATION WITH STREAMING ENGINES FOR REAL-TIME ANALYTICS



- **Directly ingest streaming data** from Amazon Kinesis Data Streams (KDS) and Amazon Managed Streaming for Apache Kafka (Amazon MSK) without staging in Amazon S3
- Uses Redshift materialized view
- **Flexible schema**: Define a schema or ingest semi-structured data with the SUPER data type
- **Speed**: Support 300K rows per second (2 KB size) with <30s latency
- **Near real-time analytics** - IOT devices, system telemetry, or clickstream data from a busy website

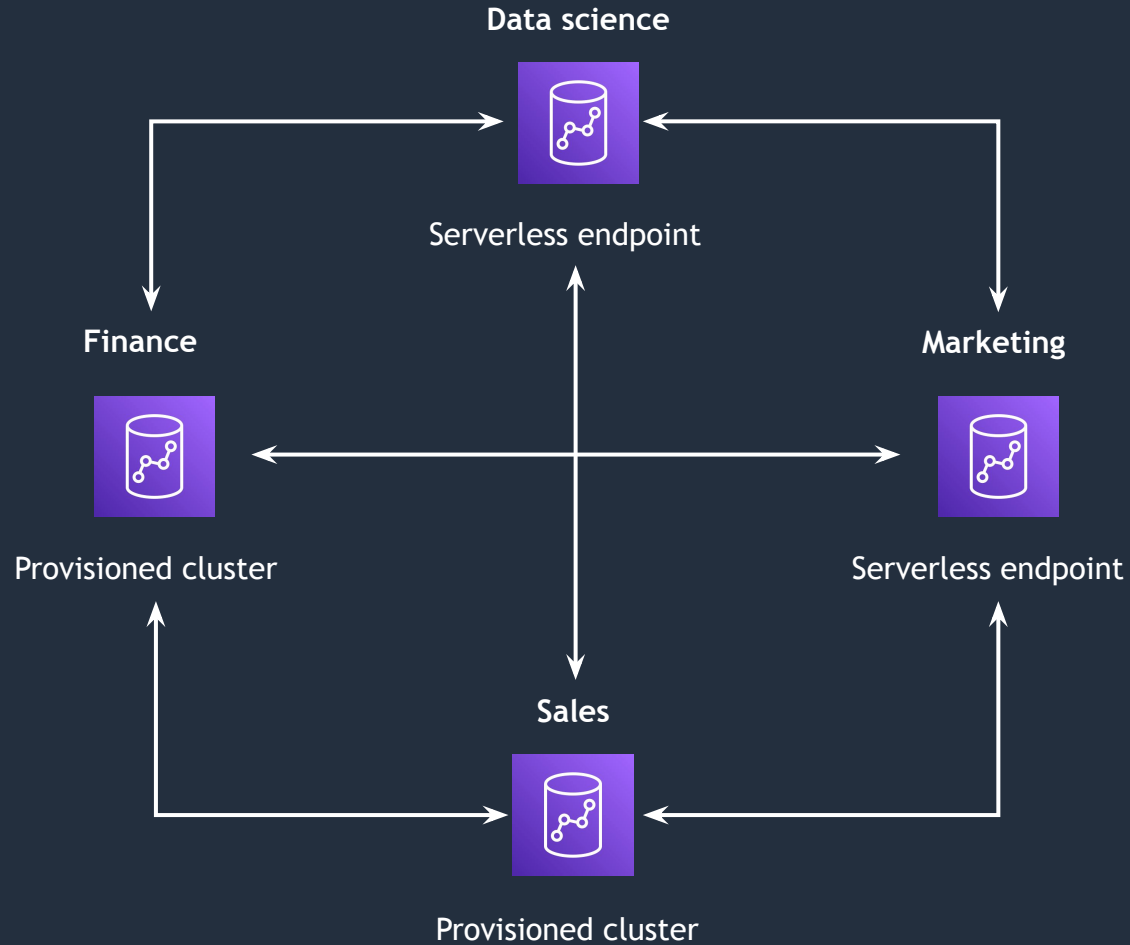
Amazon Redshift auto copy from Amazon S3



- Simplify and automate file ingestion into Amazon Redshift
- Continuously ingest data as soon as new files are created in Amazon S3
- Easily create and maintain simple data ingestion pipelines
- Automate data loading without engineering resources

Amazon Redshift Data sharing

A SECURE AND EASY WAY TO SHARE DATA ACROSS AMAZON REDSHIFT CLUSTERS



- Instant, secure & live data access without copies
- Workload independence and chargeback
- Peer-to-peer or centrally manage access to Amazon Redshift datashares with AWS Lake Formation
- Secure and governed collaboration within and across organizations (Data Mesh pattern)
- Integration with AWS Data Exchange allows you to easily find, subscribe to & use third-party data in the cloud

Built-in ML integration



Databases



Data warehouses
and data lakes




Business
intelligence tools




Amazon Aurora ML




Amazon Neptune ML



Amazon Redshift ML



Amazon Athena ML



Amazon QuickSight ML

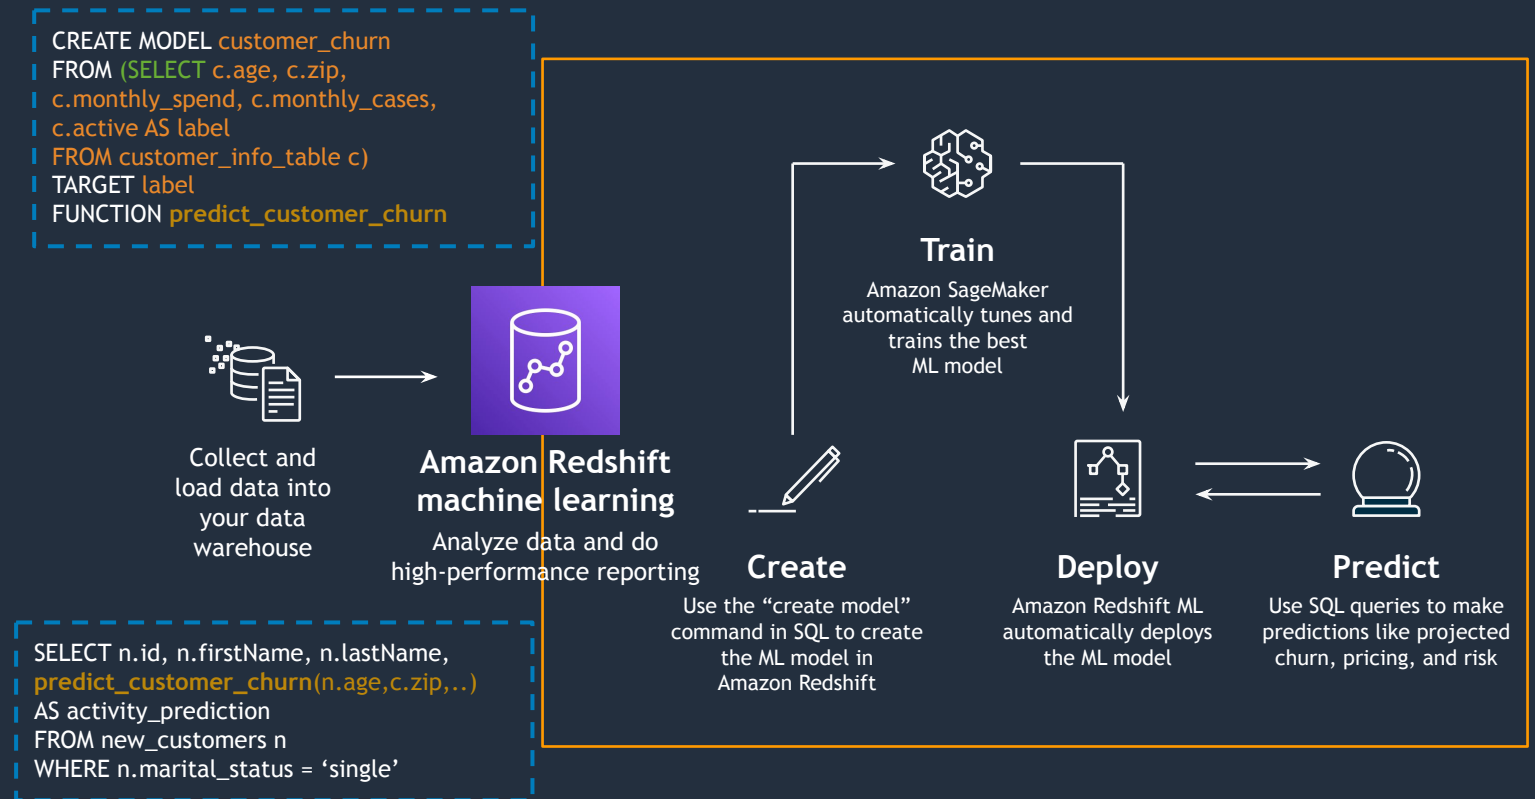
Easily create and train ML models using SQL queries with Redshift ML

Out-of-the-box integration with Amazon forecast and SageMaker for predictions/forecasting with SQL

Automatic pre-processing, creation, training, deployment, and inferencing of models

SageMaker models for in-database or remote inference

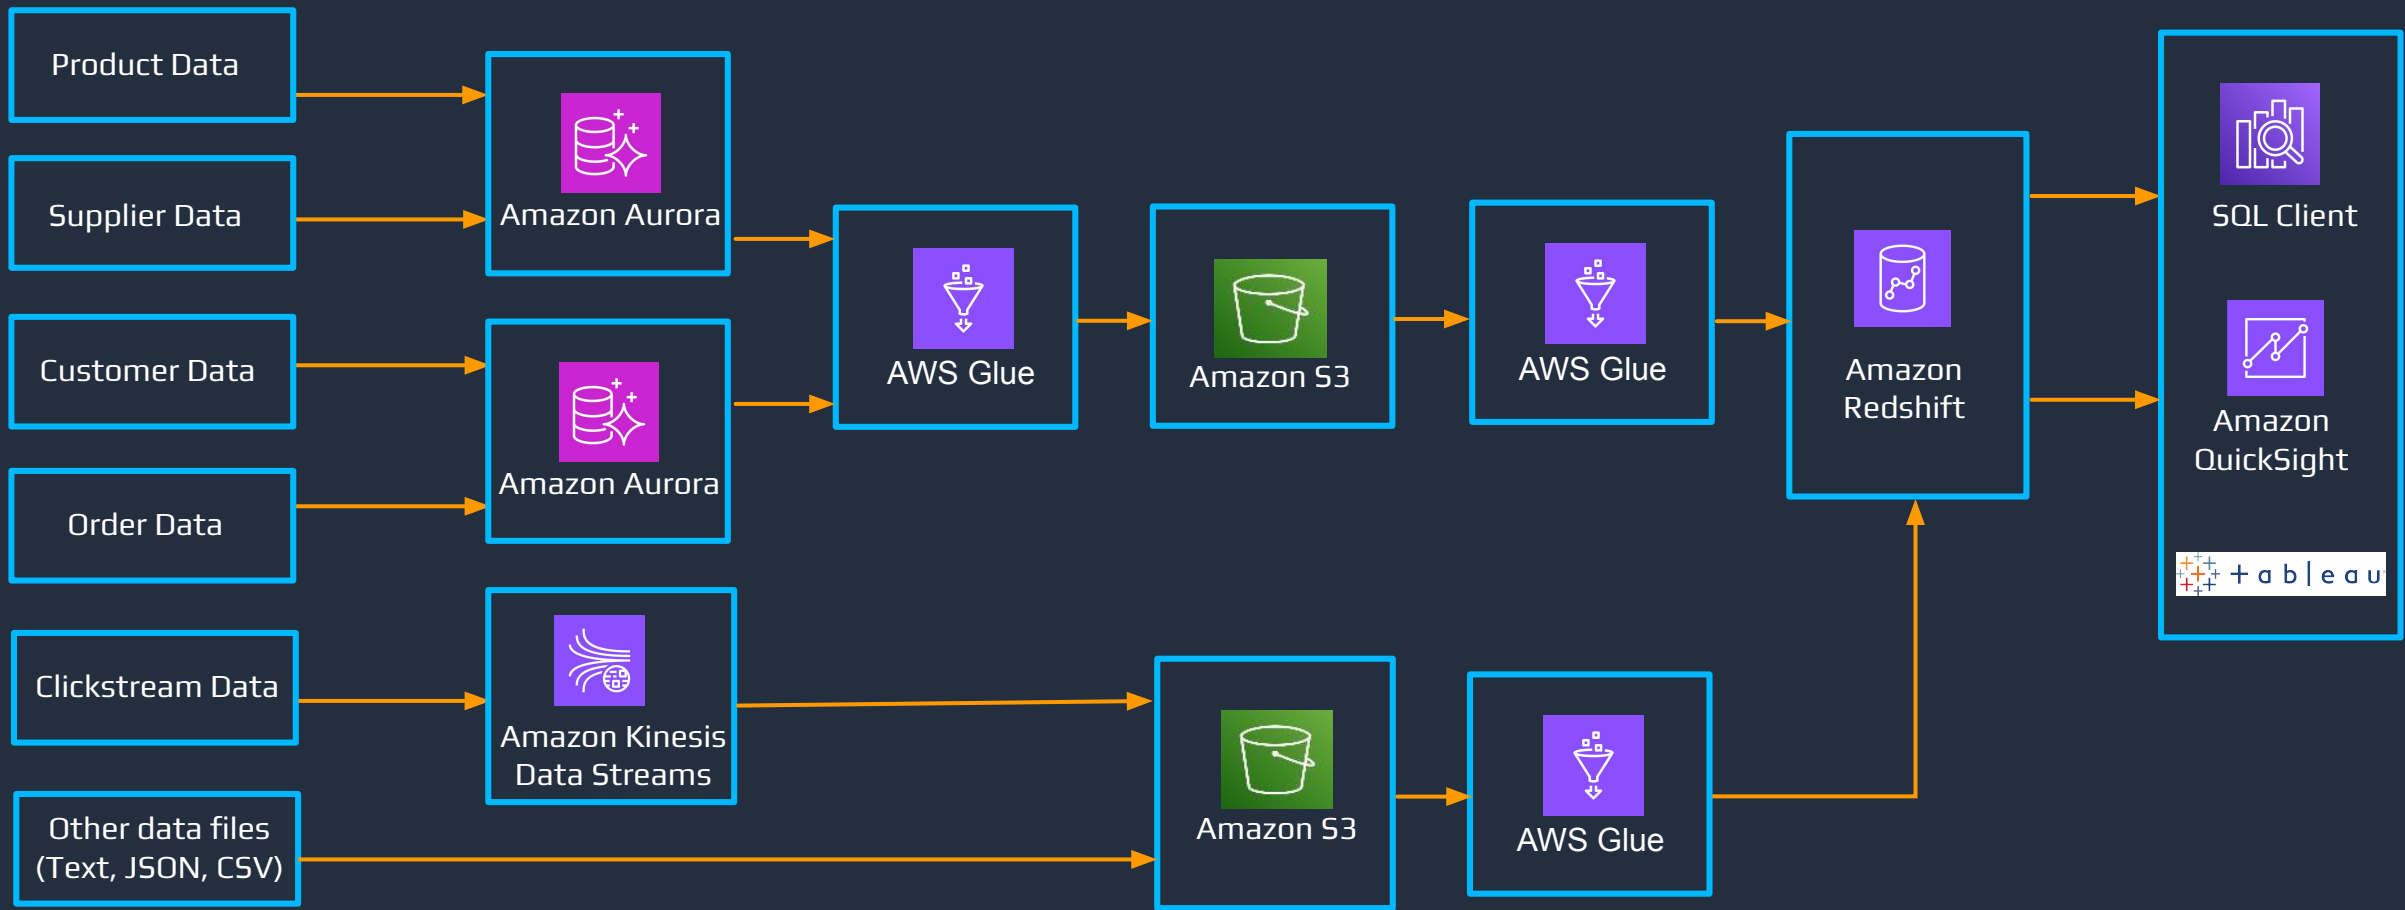
Supervised and unsupervised trainings





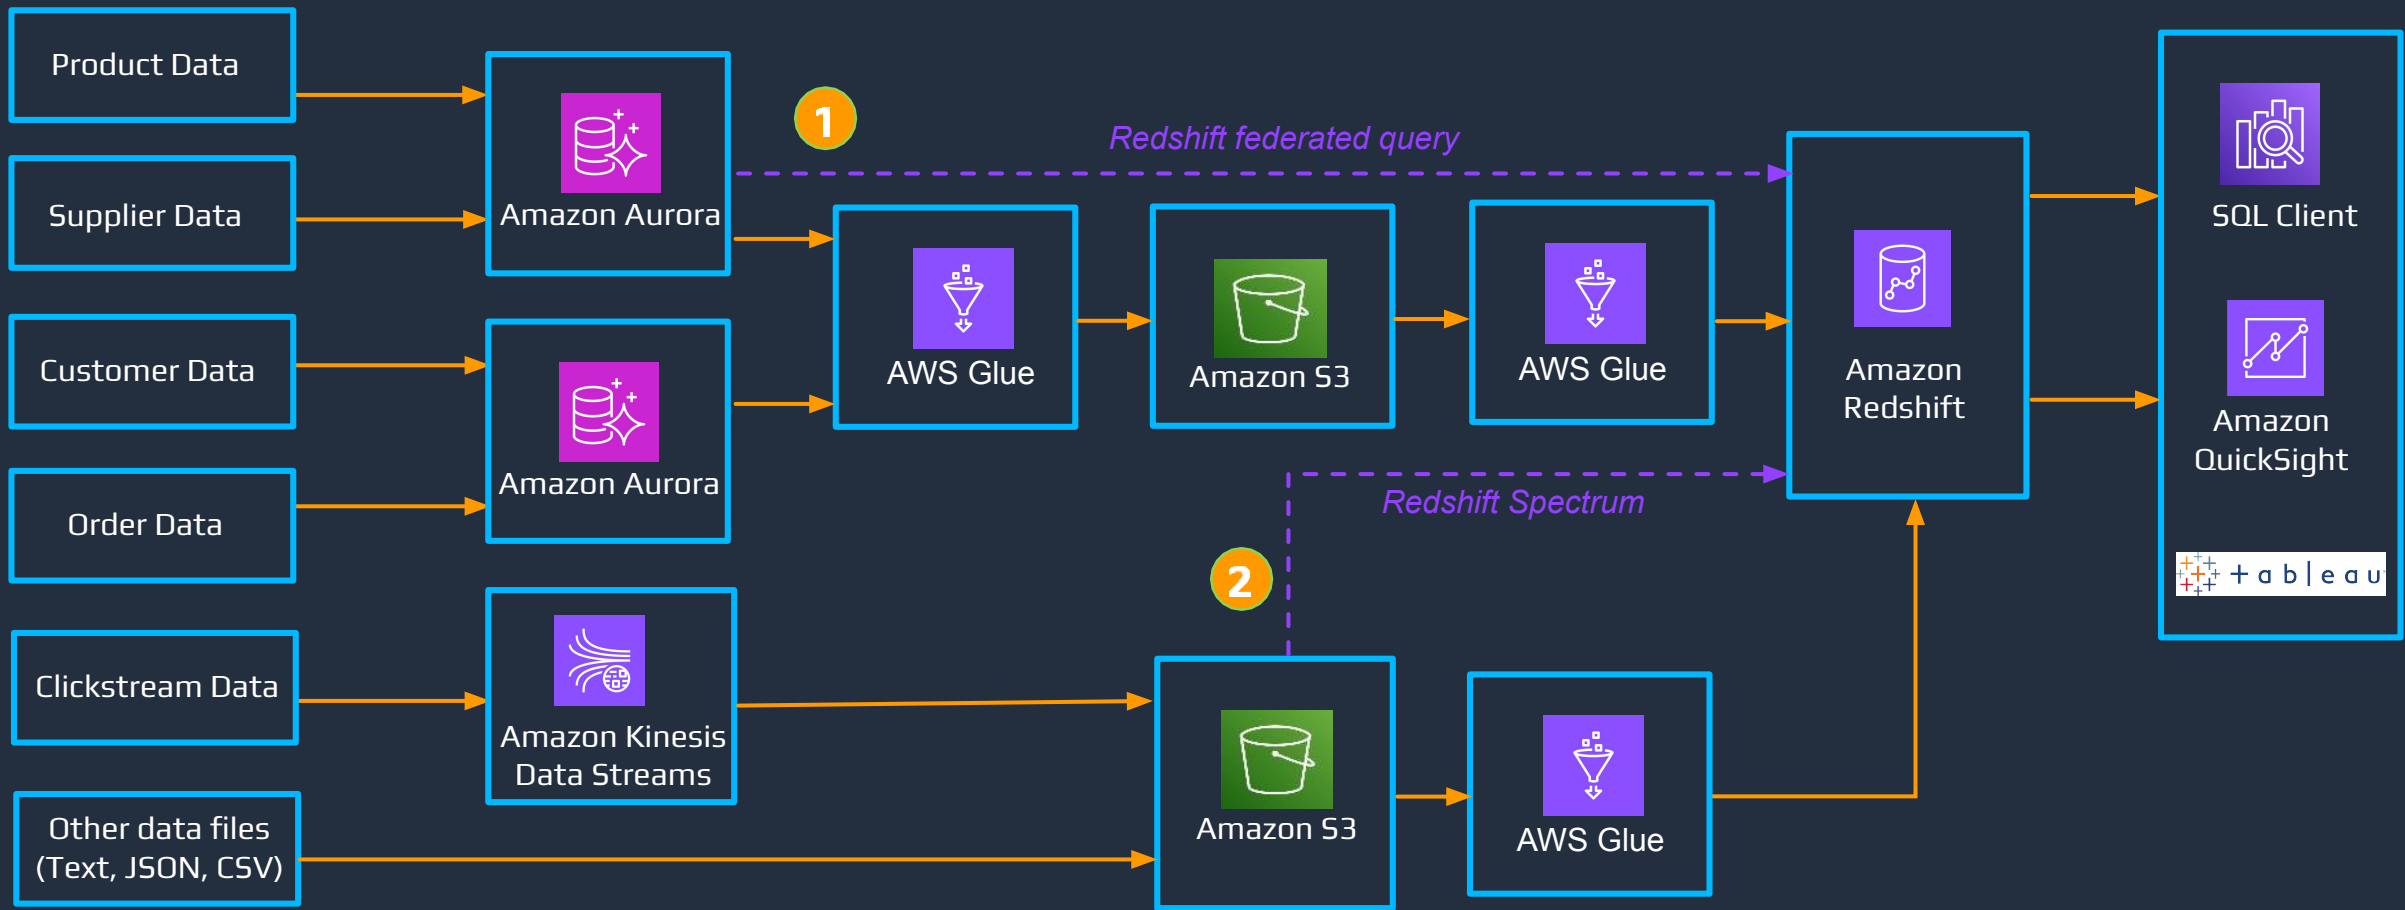
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



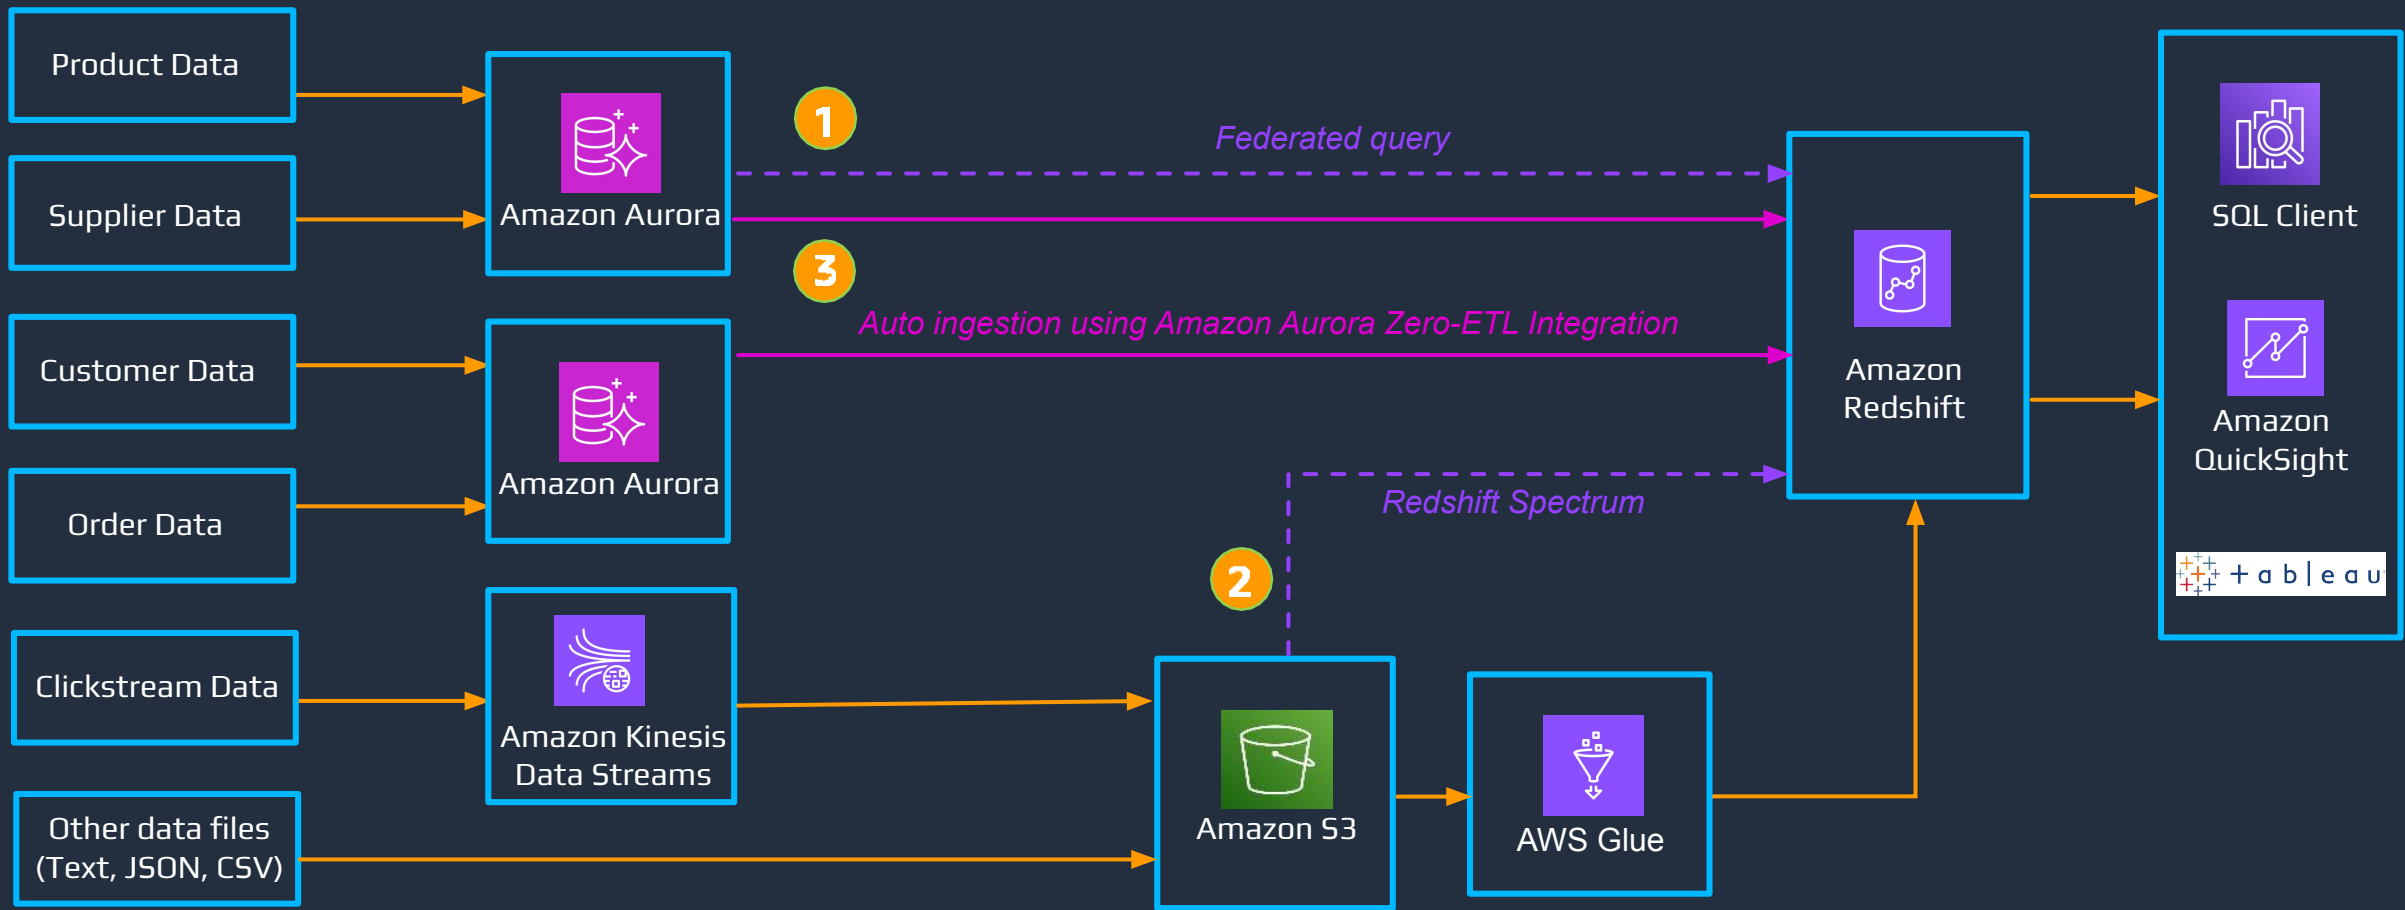
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



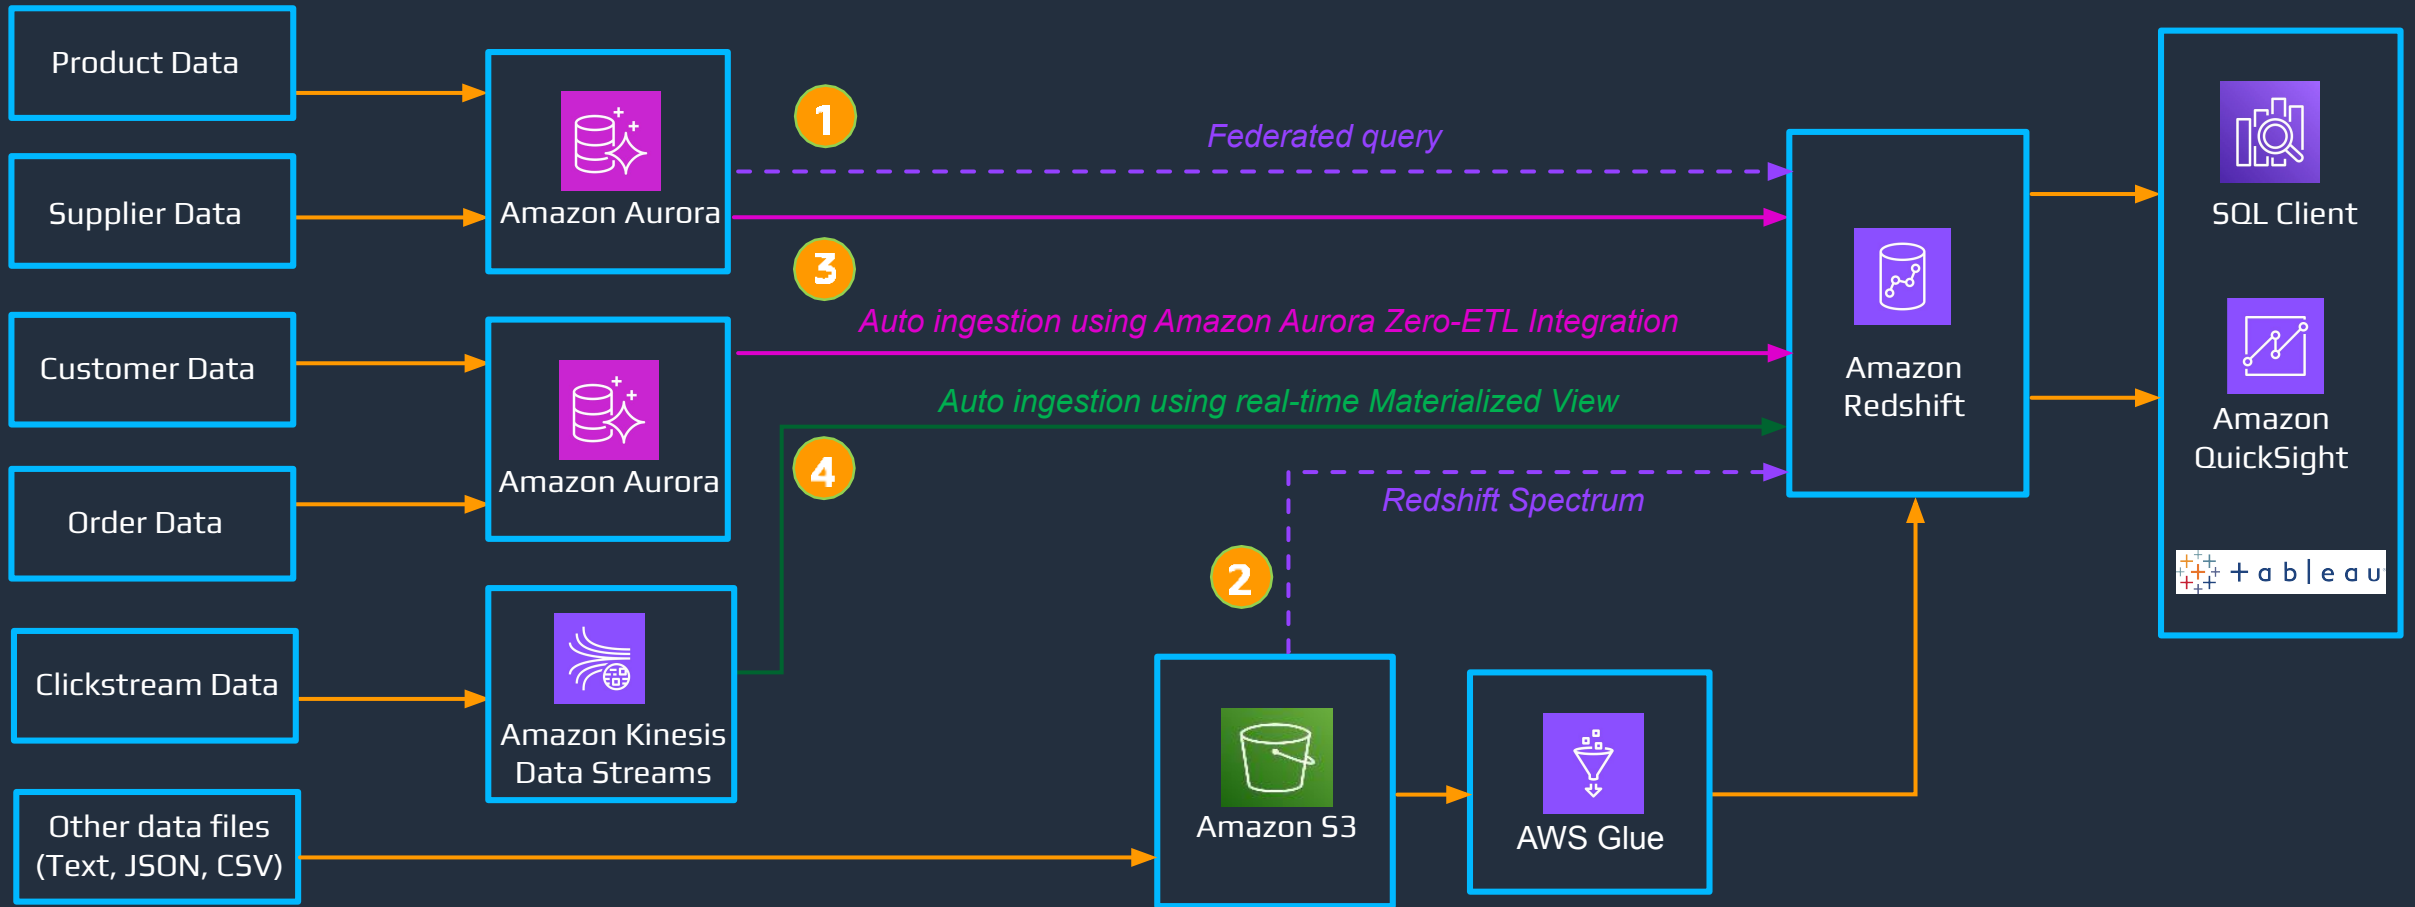
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



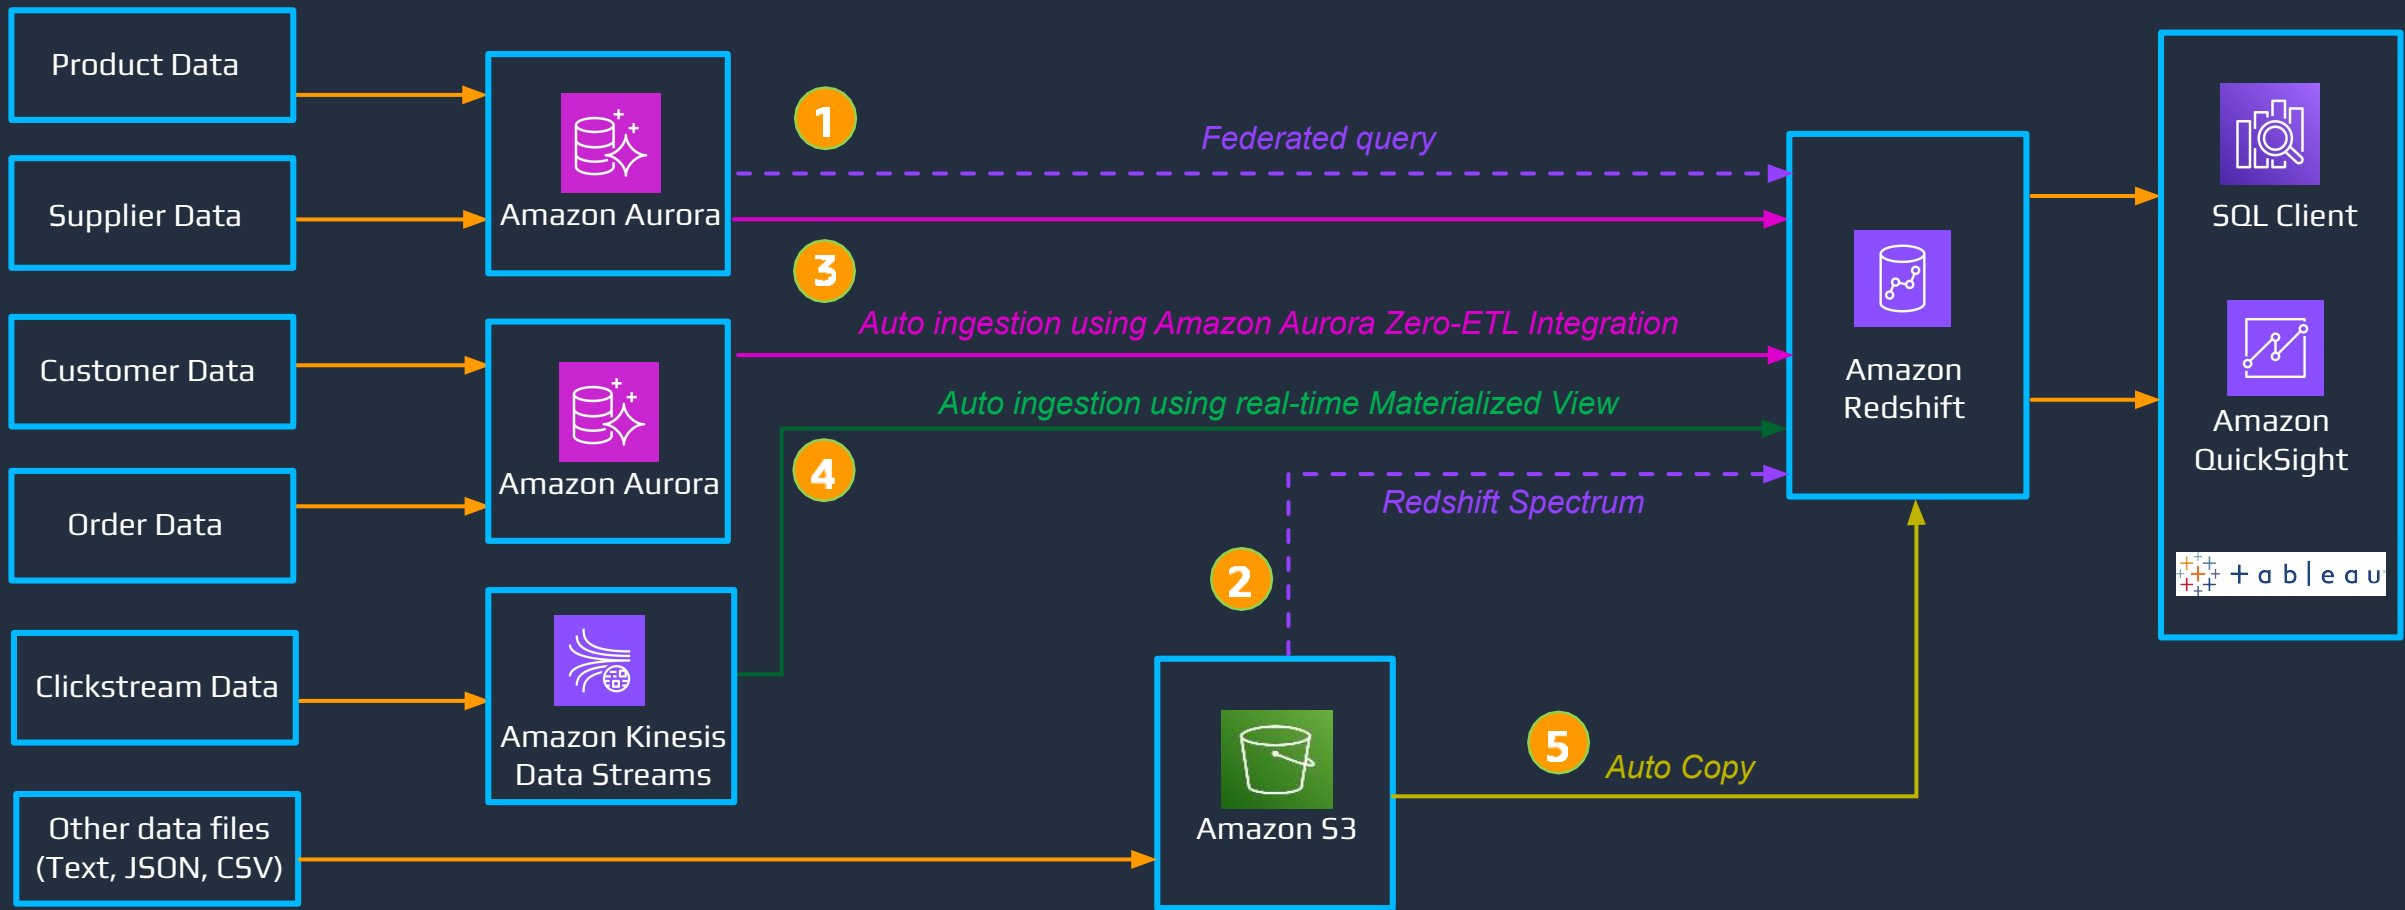
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



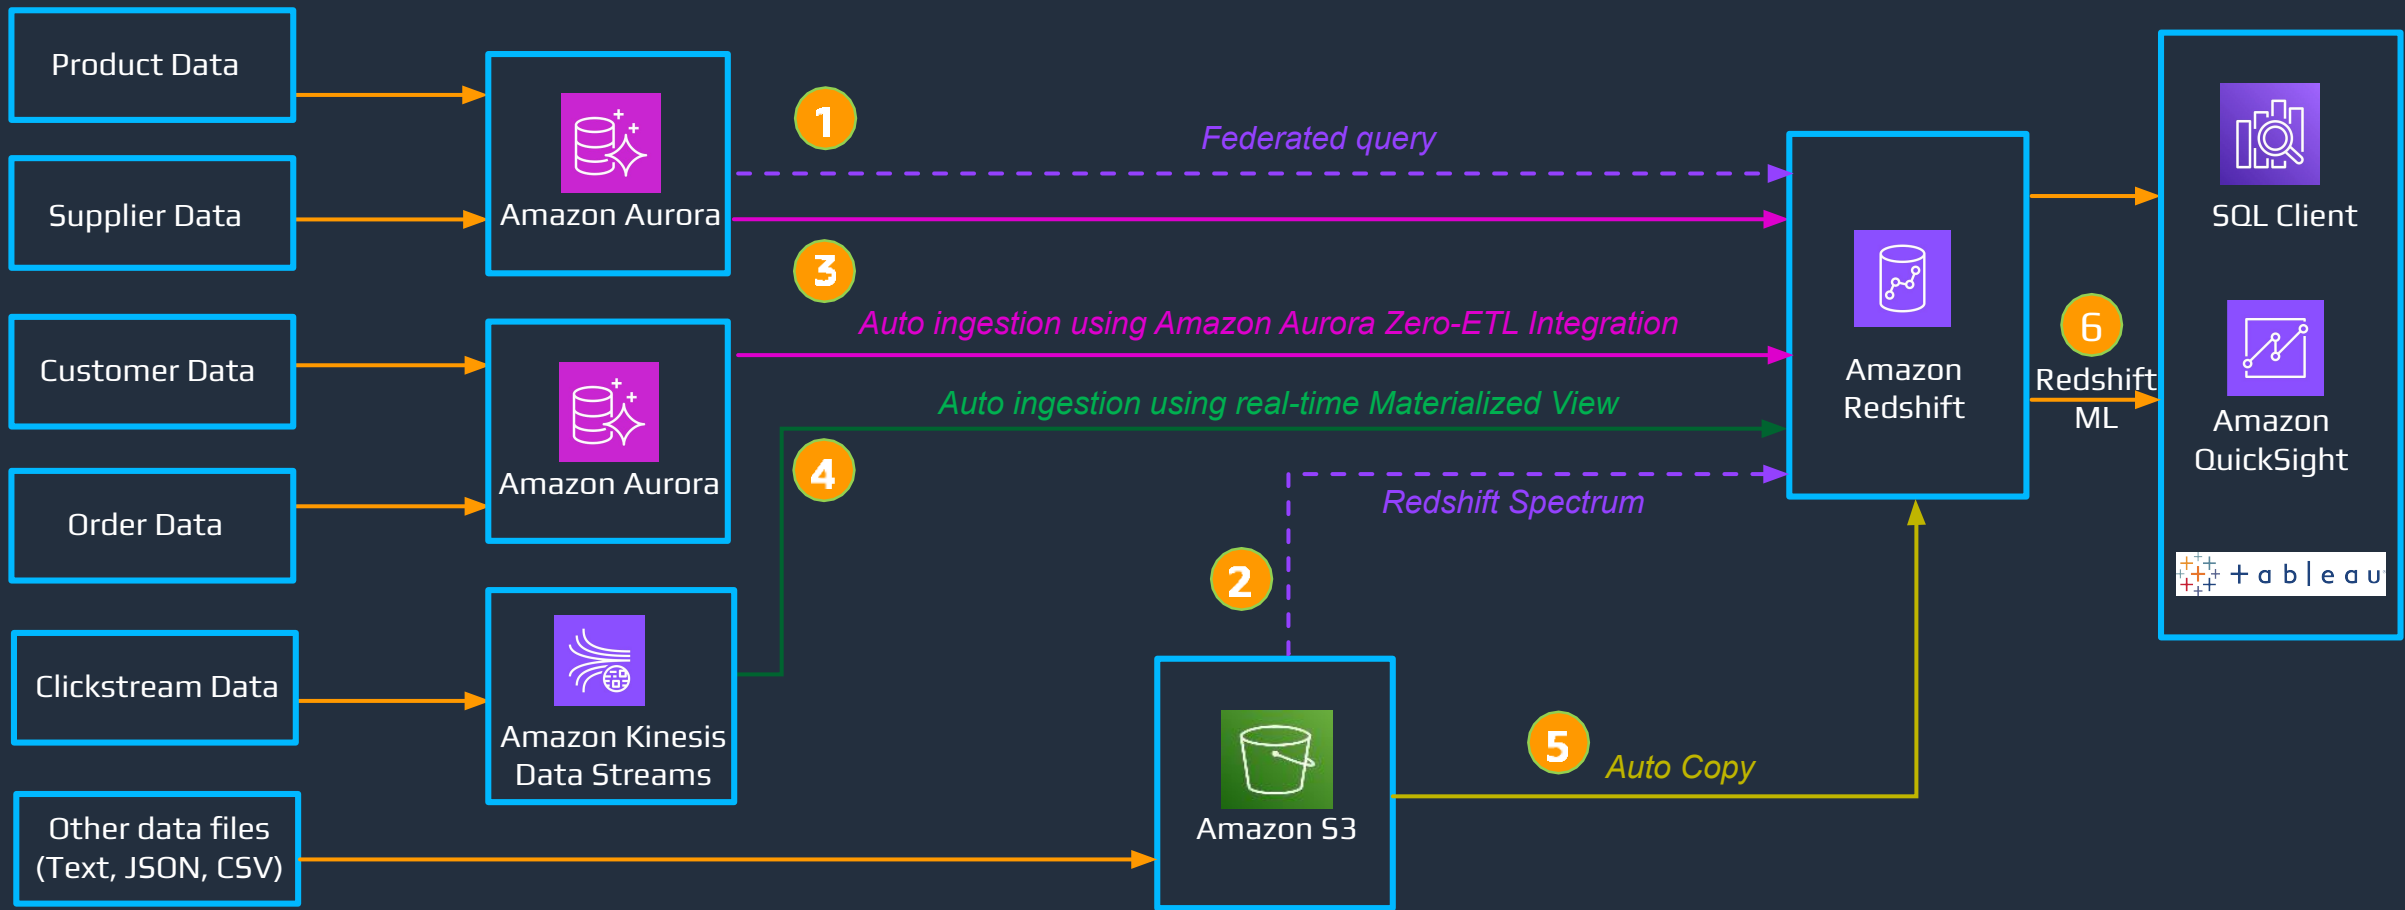
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



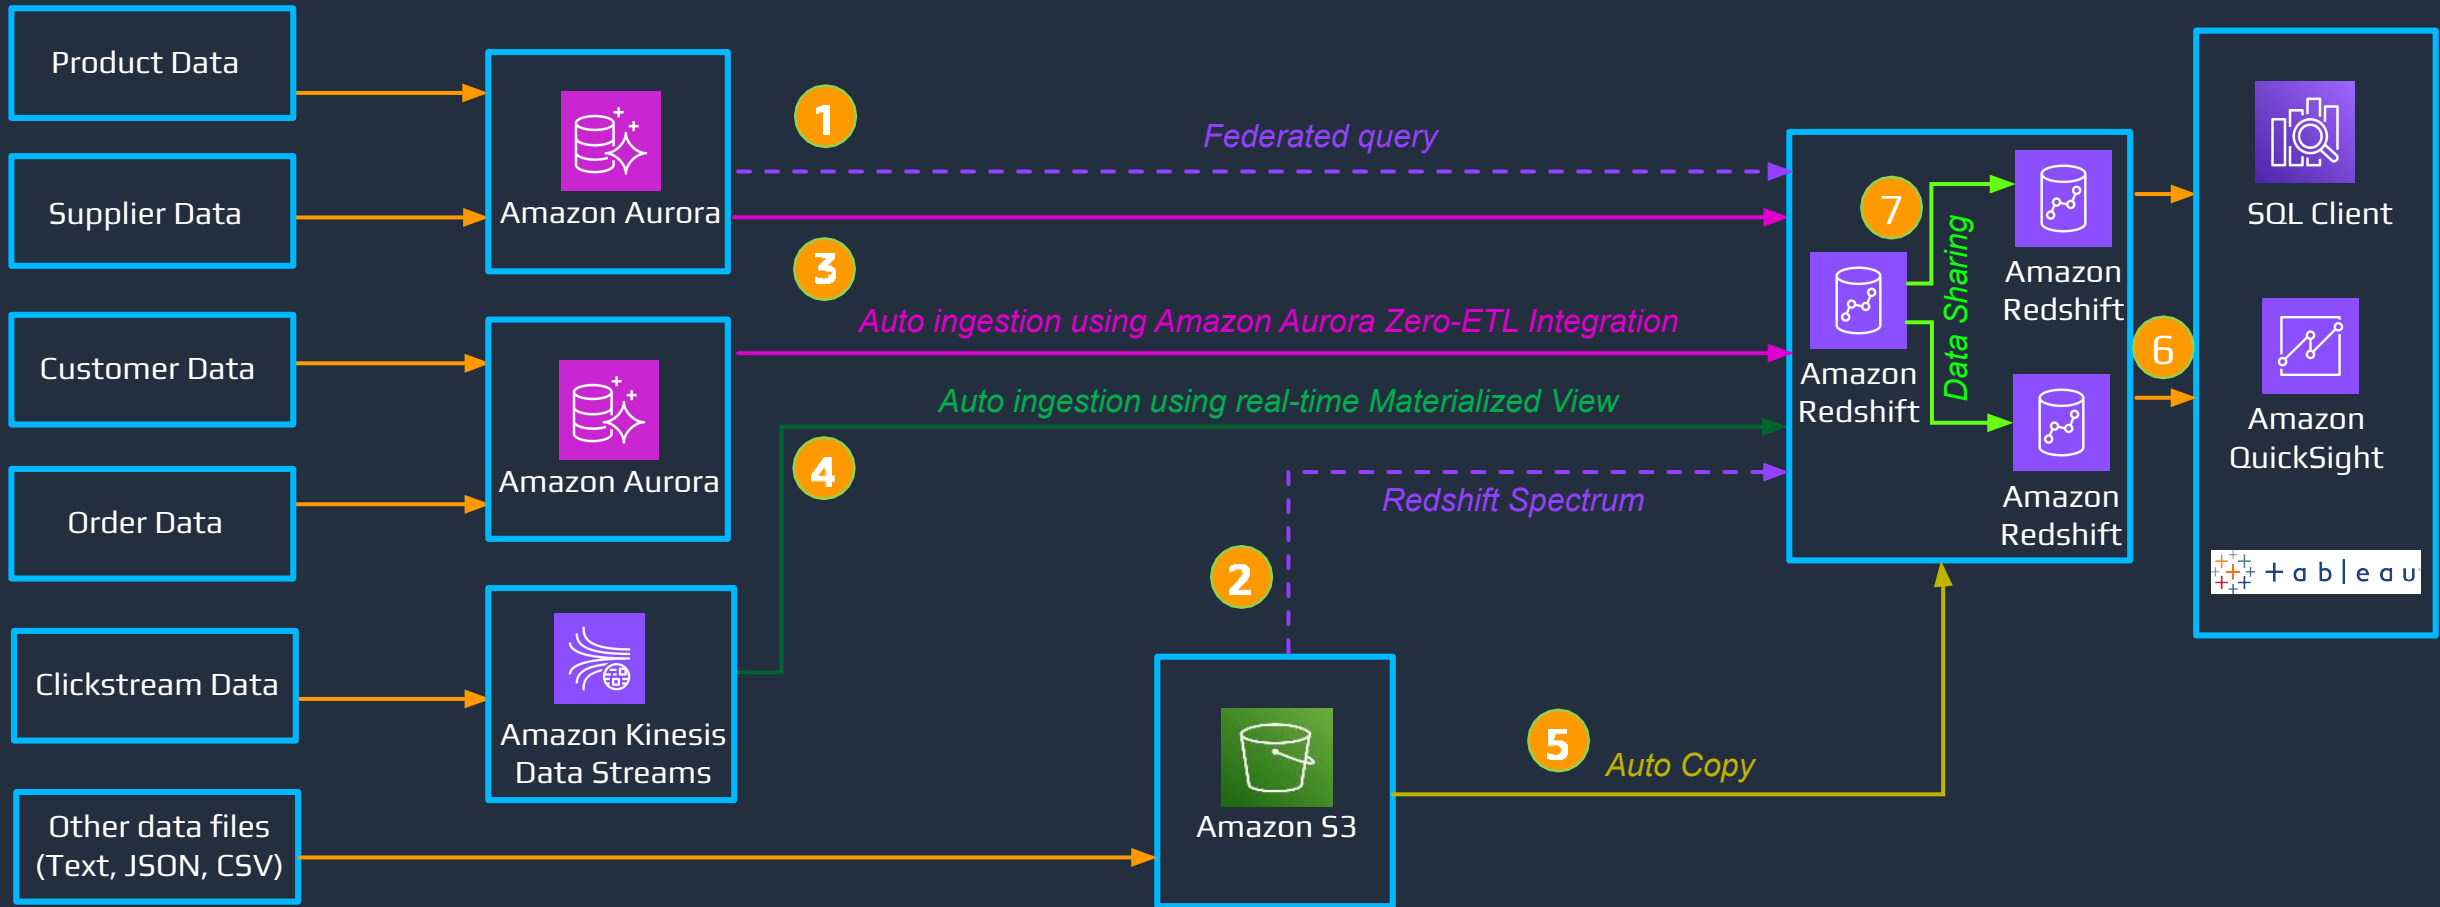
Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



Reference architecture for a Retail customer

UNIFIED REAL-TIME PROCESSING WITH REDSHIFT FEDERATED QUERY AND REDSHIFT SPECTRUM



Demo



Questions?



Thank you!

Sanchit Jain [@sanchitdilipjain](#) (LinkedIn)

